

# Deployment Logic - mmWave Radar AI System

**Assignment:** Part 4 - Deployment Logic Documentation

**Student Name:** Arnav Deshmukh

**Email:** arnavdeshmukh2004@gmail.com

**GitHub Repository:** <https://github.com/Arnav-0/GURUJI-AIR-ASSIGNMENT>

**Date:** November 2025

**Institution:** Guruji Air

## Executive Summary

This document describes the deployment architecture for a 77 GHz FMCW radar-based metal detection system using machine learning. The system achieves **82.5% classification accuracy** using Support Vector Machines (SVM) with real-time processing capability of **14.9 FPS** on CPU. This report details the complete pipeline from signal acquisition to decision output, preprocessing techniques, model inference flow, current limitations, and proposed improvements for production deployment.

### Key Achievements:

- SVM Classifier: 82.5% accuracy, 0.90 recall
- Real-time processing: 67ms per frame
- Dataset: 400 synthetic samples with balanced classes
- Preprocessing pipeline: +9% accuracy improvement in cluttered scenarios

## System Architecture

```
Hardware Radar → Signal Acquisition → Preprocessing → Model Inference → Decision Output  
(77 GHz)      (128×256 samples)      (FFT + Norm)      (SVM/CNN)      (Metal/Non-Metal)
```

## Pipeline Components

### 1. Signal Acquisition Module

- **Input:** Raw IQ data from 77 GHz FMCW radar
- **Specifications:**
  - Sample rate: 256 samples per chirp
  - Chirp rate: 128 chirps per frame
  - Frame rate: 10-30 FPS (configurable)
  - Data format: Complex float32 (I/Q channels)

### 2. Real-Time Signal Processing

```
def process_radar_frame(raw_signal):
    # Apply window function
    windowed = apply_hamming_window(raw_signal)

    # 2D FFT processing
    range_fft = fft(windowed, axis=1) # Range dimension
    doppler_fft = fft(range_fft, axis=0) # Doppler dimension

    # Generate magnitude heatmap
    heatmap = np.abs(doppler_fft)

    # Downsample to model input size (64x64)
    heatmap_resized = resize(heatmap, (64, 64))

    # Normalize
    heatmap_normalized = (heatmap_resized - mean) / std

    return heatmap_normalized
```

**Processing Time:** ~35ms per frame (CPU), ~8ms (GPU)

---

## Preprocessing Pipeline

### Stage 1: Background Subtraction

**Purpose:** Remove static clutter and environmental noise

```
def background_subtraction(current_frame, background_model):
    # Subtract learned background
    foreground = current_frame - background_model

    # Apply threshold
    foreground[foreground < threshold] = 0

    return foreground
```

**Performance Impact:** +3% accuracy improvement

### Stage 2: Noise Filtering

**Purpose:** Reduce false positives from sensor noise

#### Techniques:

- Gaussian smoothing ( $\sigma=1.0$ )
- Median filtering (kernel size:  $3 \times 3$ )
- Morphological operations (erosion + dilation)

```
def noise_filter(heatmap):
    # Gaussian blur
    smoothed = gaussian_filter(heatmap, sigma=1.0)

    # Median filter
    denoised = median_filter(smoothed, size=3)

    return denoised
```

**Performance Impact:** +6% accuracy improvement (cumulative: +9%)

### Stage 3: CFAR Detection

**Purpose:** Adaptive thresholding for target detection

**Algorithm:** Cell-Averaging CFAR (CA-CFAR)

- Guard cells: 2
- Training cells: 8
- False alarm rate:  $10^{-4}$

```
def cfar_detection(heatmap, pfa=1e-4):
    # Calculate adaptive threshold
    for each_cell in heatmap:
        guard_region = get_guard_cells(cell, size=2)
        training_region = get_training_cells(cell, size=8)

        noise_power = np.mean(training_region)
        threshold = noise_power * scaling_factor(pfa)

        if cell_value > threshold:
            detections.append(cell)

    return detections
```

## Model Flow

### Inference Pipeline

```
Input Heatmap (64x64)
↓
Preprocessing (normalization)
↓
Model Selection:
└─ SVM (Primary) → 82.5% accuracy, <5ms inference
└─ CNN (Backup) → 51.2% accuracy, ~40ms inference
↓
```

```

Confidence Thresholding (>0.7)
↓
Post-processing (smoothing over 5 frames)
↓
Final Decision: Metal / Non-Metal

```

## Model Selection Logic

### Primary: SVM Classifier

- **Advantages:**

- Fast inference (<5ms on CPU)
- High accuracy (82.5%)
- Excellent recall (0.90) - fewer false negatives
- Low memory footprint (~100KB)

- **Feature Extraction:**

- Flatten  $64 \times 64$  heatmap → 4096 features
- Apply PCA for dimensionality reduction (optional)
- Normalize features

### Backup: CNN Classifier

- **Use Case:** When computational resources allow
- **Architecture:** 2.27M parameters
- **Inference Time:** 40ms (CPU), 8ms (GPU)

## Deployment Code

```

class RadarClassifier:
    def __init__(self):
        self.svm = load_model('svm_model.pkl')
        self.scaler = load_scaler('scaler.pkl')
        self.frame_buffer = []

    def predict(self, heatmap):
        # Flatten and normalize
        features = heatmap.flatten()
        features_scaled = self.scaler.transform(features)

        # SVM prediction
        prediction = self.svm.predict(features_scaled)
        confidence = self.svm.predict_proba(features_scaled)

        # Temporal smoothing
        self.frame_buffer.append(prediction)
        if len(self.frame_buffer) > 5:
            self.frame_buffer.pop(0)

```

```

# Majority voting
final_prediction = mode(self.frame_buffer)

return final_prediction, confidence

```

## Real-Time Performance

<b>Component</b>	<b>Time (ms)</b>	<b>Throughput</b>
Signal Acquisition	10	100 FPS
2D FFT Processing	35	28.6 FPS
Preprocessing	15	66.7 FPS
SVM Inference	5	200 FPS
Post-processing	2	500 FPS
<b>Total Pipeline</b>	<b>67ms</b>	<b>14.9 FPS</b>

**Note:** With GPU acceleration, total time reduces to ~35ms (28.6 FPS)

---

## Limitations and Improvements

### Current Limitations

#### 1. Dataset Limitations

- **Synthetic Data Only:** All training data is simulated
- **Limited Scenarios:** Only binary classification (metal/non-metal)
- **Simplified Physics:** Real-world radar has multipath, interference
- **Impact:** Model may not generalize to real hardware

**Evidence:** CNN performance (51.2%) suggests overfitting to synthetic patterns

#### 2. Model Performance

- **CNN Underperformance:** 51.2% accuracy (barely better than random)
- **Root Causes:**
  - Learning rate issues
  - Insufficient training epochs
  - Possible class imbalance in training
- **Impact:** Cannot reliably use CNN as backup model

#### 3. Real-Time Constraints

- **CPU Bottleneck:** 67ms per frame limits to 14.9 FPS
- **Latency:** Total system latency ~100ms including decision logic
- **Impact:** May miss fast-moving objects or require frame skipping

## 4. Environmental Sensitivity

- **Weather:** Rain, fog affect radar signal quality
- **Clutter:** Dense environments create false positives
- **Range:** Performance degrades beyond 5 meters
- **Impact:** Detection accuracy drops from 82.5% to ~64% in clutter

## 5. Hardware Dependencies

- **Calibration Required:** Each radar unit needs individual calibration
- **Temperature Drift:** Sensor performance varies with temperature
- **Power Consumption:** Continuous operation requires ~5W power
- **Impact:** Field deployment requires careful integration

## Proposed Improvements

### Short-Term (1-3 months)

#### 1. Real-World Data Collection

- **Action:** Deploy prototype with labeled data collection
- **Goal:** 1000+ real samples from actual radar hardware
- **Expected Impact:** +15-20% accuracy improvement

#### 2. CNN Architecture Optimization

- **Actions:**
  - Reduce learning rate ( $0.001 \rightarrow 0.0001$ )
  - Increase training epochs ( $50 \rightarrow 100$ )
  - Add more aggressive data augmentation
  - Implement focal loss for class imbalance
- **Expected Impact:** CNN accuracy  $51.2\% \rightarrow 75\%-80\%$

#### 3. GPU Acceleration

- **Action:** Implement CUDA/TensorRT inference
- **Hardware:** NVIDIA Jetson Nano or better
- **Expected Impact:**  $67\text{ms} \rightarrow 35\text{ms}$  (28.6 FPS real-time)

### Medium-Term (3-6 months)

## 4. Multi-Class Classification

- **Extension:** Beyond metal/non-metal
  - Plastic
  - Wood
  - Glass
  - Composite materials
- **Dataset Required:** 500+ samples per class
- **Expected Accuracy:** 70-75% (5-class)

## 5. Advanced Preprocessing

- **Techniques:**
  - Adaptive CFAR with multiple thresholds
  - Machine learning-based clutter removal
  - Temporal filtering across multiple frames
- **Expected Impact:** +5-8% accuracy in cluttered environments

## 6. Ensemble Methods

- **Approach:** Combine SVM + CNN + Random Forest
- **Voting Scheme:** Weighted by individual confidence
- **Expected Impact:** +3-5% accuracy improvement

### Long-Term (6-12 months)

## 7. Deep Learning Enhancements

- **Architecture:** Attention mechanisms (Transformer-based)
- **Benefits:** Better feature extraction, spatial relationships
- **Expected Accuracy:** 85-90%

## 8. Edge Deployment Optimization

- **Techniques:**
  - Model quantization (FP32 → INT8)
  - Knowledge distillation (compress 2.27M → 500K params)
  - Pruning (remove 40-50% of weights)
- **Impact:** 2-3x faster inference, 70% less memory

## 9. Multi-Sensor Fusion

- **Integration:** Combine radar with camera/LiDAR
- **Benefits:** Redundancy, cross-validation
- **Expected Accuracy:** 90-95%

## 10. Adaptive Learning

- **Approach:** Online learning with user feedback
- **Benefits:** Continuously improve in deployment
- **Implementation:** Federated learning across multiple units

---

## Deployment Checklist

### Pre-Deployment

- Collect 500+ real-world samples for validation
- Calibrate radar hardware in target environment
- Benchmark inference speed on deployment hardware
- Test in various environmental conditions (rain, fog, night)

- Establish baseline accuracy metrics

## Integration

- Implement data acquisition interface
- Set up preprocessing pipeline
- Load trained models (SVM primary, CNN backup)
- Configure confidence thresholds
- Implement logging and monitoring

## Validation

- A/B testing against manual inspection
- Measure false positive/negative rates
- Test latency under load
- Verify power consumption
- Document failure modes

## Monitoring

- Real-time accuracy tracking
- Alert on confidence drops
- Log edge cases for retraining
- Track model drift over time
- Schedule periodic recalibration

---

## Conclusion & Future Roadmap

The mmWave radar AI system demonstrates strong feasibility for automated metal detection with **82.5% classification accuracy** using SVM. The real-time pipeline achieves **14.9 FPS on CPU**, which is adequate for many industrial and security screening applications. The preprocessing pipeline successfully improves detection accuracy by **9%** in cluttered environments through background subtraction and noise filtering techniques.

## Critical Next Steps

### Immediate Priorities (1-3 months):

1. **Real-world data collection:** Validate synthetic training with 1000+ samples from actual radar hardware
2. **GPU acceleration:** Reduce pipeline latency from 67ms to 35ms for 28.6 FPS real-time performance
3. **CNN optimization:** Improve CNN accuracy from 51.2% to 75-80% through hyperparameter tuning and architectural changes
4. **Field testing:** Deploy prototype in target environments to measure real-world performance degradation

### Medium-term Goals (3-6 months):

- Expand to multi-class classification (metal, plastic, wood, glass, composite)
- Implement ensemble methods combining SVM + CNN + Random Forest for improved accuracy

- Develop adaptive preprocessing for varying environmental conditions
- Optimize for edge deployment with model quantization and pruning

### Long-term Vision (6-12 months):

- Achieve 90%+ accuracy with deep learning enhancements (attention mechanisms, transformers)
- Multi-sensor fusion integrating radar with camera/LiDAR for redundancy
- Online learning capability with federated learning across multiple deployed units
- Full edge deployment on embedded systems (NVIDIA Jetson, mobile devices)

## Production Readiness Assessment

**Current Status:** Proof-of-concept validated

Criterion	Notes
Algorithm Performance	SVM meets 82.5% requirement
Real-time Processing	14.9 FPS adequate but GPU recommended
Robustness	Only synthetic data validated
Hardware Integration	Requires actual radar interfacing
Safety/Security	Needs dual-validation and human-in-loop
Scalability	Modular architecture supports scaling

**Recommendation:** System is ready for **controlled pilot deployment** with human oversight. Not recommended for fully autonomous operation until real-world validation is complete.

## Key Takeaways

1. **SVM outperforms CNN** on this dataset due to limited training samples (400) relative to CNN complexity (2.27M parameters)
2. **Preprocessing is critical:** Simple techniques (background subtraction, noise filtering) provide significant accuracy gains (+9%)
3. **Real-time is achievable:** 67ms latency on CPU is acceptable for non-safety-critical applications; GPU can reduce to 35ms
4. **Synthetic-to-real gap** is the primary deployment risk that must be addressed through field data collection
5. **Modular design** enables incremental improvements without system redesign

## References & Resources

### Technical Documentation:

- Project README: [GitHub Repository](#)
- Jupyter Notebooks: [/notebooks/](#) (complete implementation)
- Source Code: [/src/](#) (radar simulator, models, preprocessing)
- Results: [/outputs/](#) (27 visualizations, trained models, metrics)

**Performance Metrics:**

- Classification Results: [/outputs/results/classification\\_results.json](#)
- Detection Results: [/outputs/results/hidden\\_object\\_detection\\_results.json](#)
- Trained Models: [/data/models/](#) (SVM, CNN checkpoints)

**Contact Information:**

- **Author:** Arnav Deshmukh
  - **Email:** arnavdeshmukh2004@gmail.com
  - **GitHub:** <https://github.com/Arnav-0/GURUJI-AIR-ASSIGNMENT>
- 

**Document Version:** 1.0

**Last Updated:** November 22, 2025

**Project:** mmWave Radar AI - Deployment Logic Documentation

**Assignment:** Guruji Air - Part 4

---

*This document is part of the mmWave Radar AI assignment submission. For complete implementation details, please refer to the GitHub repository and accompanying Jupyter notebooks.*